

## Lab 2

1. Load the fetch\_20newsgroups dataset and split it into training and test subsets.

```
from sklearn.datasets import fetch_20newsgroups
import nltk
from nltk.corpus import stopwords
from sklearn.feature_extraction.text import CountVectorizer,
TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score,
recall_score, confusion_matrix
import matplotlib.pyplot as plt

nltk.download('stopwords')
nltk.download('punkt')

newsgroups_train = fetch_20newsgroups(subset='train')
newsgroups_test = fetch_20newsgroups(subset='test')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
```

## 2. Text preprocessing

```
def preprocess_texts(texts : list[str]) -> list[str]:
    """
    Remove all stopwords, remove non-letter words, lowercase.
    """
    filtered = []
    for text in texts:
        temp = []
        for word in nltk.tokenize.word_tokenize(text):
            if word.lower() not in stopwords.words('english') and
word.isalpha():
                temp.append(word.lower())

        filtered.append(" ".join(temp))
    return filtered

def stem_texts(texts : list[str]):
```

```

'''
Stem all words in texts.\n\nAccepts already tokenized texts.
'''

stemmed = []
stemmer = nltk.SnowballStemmer("english")

for text in texts:
    temp = []
    for word in nltk.tokenize.word_tokenize(text):
        temp.append(stemmer.stem(word))

    stemmed.append(" ".join(temp))
return stemmed

newsgroups_train['preprocessed_data'] =
preprocess_texts(newsgroups_train.data)
newsgroups_test['preprocessed_data'] =
preprocess_texts(newsgroups_test.data)

newsgroups_train['stemmed_data'] =
stem_texts(newsgroups_train['preprocessed_data'])
newsgroups_test['stemmed_data'] =
stem_texts(newsgroups_test['preprocessed_data'])

```

### 3. Vectorizers

```

def vectorize_data(vectorizer, train_data, test_data):
'''
Fits vectorizer and vectorizes data.

Returns: Vectorized train data, vectorized test data.
'''

vec_train = vectorizer.fit_transform(train_data)
vec_test = vectorizer.transform(test_data)
return vec_train, vec_test

'''
Count Vectorizer
'''

# raw texts
raw_text_count_train, raw_text_count_test = vectorize_data(
    CountVectorizer(),
    newsgroups_train.data,
    newsgroups_test.data)

# preprocessed texts excluding stemming
preprocessed_count_train, preprocessed_count_test = vectorize_data(

```

```

    CountVectorizer(),
    newsgroups_train['preprocessed_data'],
    newsgroups_test['preprocessed_data'])

# preprocessed texts including stemming
stemmed_count_train, stemmed_count_test = vectorize_data(
    CountVectorizer(),
    newsgroups_train['stemmed_data'],
    newsgroups_test['stemmed_data'])

# preprocessed texts considering bigrams as well
stemmed_bigram_count_train, stemmed_bigram_count_test =
vectorize_data(
    CountVectorizer(ngram_range=(1,2)),
    newsgroups_train['stemmed_data'],
    newsgroups_test['stemmed_data'])

'''
Tf-idf Vectorizer
'''

# raw texts
raw_text_tfidf_train, raw_text_tfidf_test = vectorize_data(
    TfidfVectorizer(),
    newsgroups_train.data,
    newsgroups_test.data)

# preprocessed texts excluding stemming
preprocessed_tfidf_train, preprocessed_tfidf_test = vectorize_data(
    TfidfVectorizer(),
    newsgroups_train['preprocessed_data'],
    newsgroups_test['preprocessed_data'])

# preprocessed texts including stemming
stemmed_tfidf_train, stemmed_tfidf_test = vectorize_data(
    TfidfVectorizer(),
    newsgroups_train['stemmed_data'],
    newsgroups_test['stemmed_data'])

# preprocessed texts considering bigrams as well
stemmed_bigram_tfidf_train, stemmed_bigram_tfidf_test =
vectorize_data(
    TfidfVectorizer(ngram_range=(1,2)),
    newsgroups_train['stemmed_data'],
    newsgroups_test['stemmed_data'])

```

## 4. Questions

How many instances do training and test subsets of the fetch\_20newsgroups dataset contain separately?

```
print("The training subset has ", len(newsgroups_train.data), "  
instances.")  
print("The test subset has ", len(newsgroups_test.data), "  
instances.")
```

```
The training subset has 11314 instances.  
The test subset has 7532 instances.
```

How many classes are there for the articles in the dataset, and what are the names of these classes?

```
print("There are ", len(newsgroups_test.target_names), " different  
classes.")
```

```
for name in newsgroups_test.target_names:  
    print(name)
```

```
There are 20 different classes.
```

```
alt.atheism  
comp.graphics  
comp.os.ms-windows.misc  
comp.sys.ibm.pc.hardware  
comp.sys.mac.hardware  
comp.windows.x  
misc.forsale  
rec.autos  
rec.motorcycles  
rec.sport.baseball  
rec.sport.hockey  
sci.crypt  
sci.electronics  
sci.med  
sci.space  
soc.religion.christian  
talk.politics.guns  
talk.politics.mideast  
talk.politics.misc  
talk.religion.misc
```

How does the dimension of the CountVector vary based on different preprocessing approaches?

```
print("Dimension of count vector for raw text: ",  
raw_text_count_train.shape[1])
```

```

print("Dimension of count vector for preprocessed data: ",
preprocessed_count_train.shape[1])
print("Dimension of count vector for stemmed data: ",
stemmed_count_train.shape[1])
print("Dimension of count vector for stemmed data considering bigrams:
", stemmed_bigram_count_train.shape[1])

```

```

Dimension of count vector for raw text: 130107
Dimension of count vector for preprocessed data: 73078
Dimension of count vector for stemmed data: 54944
Dimension of count vector for stemmed data considering bigrams:
867347

```

The more preprocessed the text is, the less distinct tokens are there, the smaller the dimension of the vectors.

When considering bigrams, dimensionality goes up because of the number of combinations.

## 5. Classify the texts

```

def get_predictions_of_model_on_data(model_class, train_data,
test_data, train_target):
    """
    Fit a model with vectorized data and get predictions.
    """
    model = model_class.fit(train_data, train_target)
    return model.predict(test_data)

```

Train each model on each type of preprocessing and vectorizer, let the model classify the test data and save the predictions in a `results_dict`

```

results_dict = {}
for classifier, name in [(MultinomialNB(), "Naive Bayes"), (SVC(),
"SVC"), (LogisticRegression(), "Regression"),
(RandomForestClassifier(), "Random Forest")]:
    print("Training", name)
    temp = {}

    # raw texts
    temp["raw texts count"] =
get_predictions_of_model_on_data(classifier,
raw_text_count_train,
raw_text_count_test,
newsgroups_train.target)

    temp["raw texts tfidf"] =
get_predictions_of_model_on_data(classifier,
raw_text_tfidf_train,
raw_text_tfidf_test,

```

```

newsgroups_train.target)

    # preprocessed texts excluding stemming
    temp["preprocessed texts count"] =
get_predictions_of_model_on_data(classifier,
preprocessed_count_train,
preprocessed_count_test,
newsgroups_train.target)

    temp["preprocessed texts tfidf"] =
get_predictions_of_model_on_data(classifier,
preprocessed_tfidf_train,
preprocessed_tfidf_test,
newsgroups_train.target)

    # preprocessed texts including stemming
    temp["stemmed texts count"] =
get_predictions_of_model_on_data(classifier,
                                stemmed_count_train,
                                stemmed_count_test,
newsgroups_train.target)

    temp["stemmed texts tfidf"] =
get_predictions_of_model_on_data(classifier,
                                stemmed_tfidf_train,
                                stemmed_tfidf_test,
newsgroups_train.target)

    # preprocessed texts considering bigrams as well
    temp["stemmed bigram texts count"] =
get_predictions_of_model_on_data(classifier,
stemmed_bigram_count_train,
stemmed_bigram_count_test,
newsgroups_train.target)

    temp["stemmed bigram texts tfidf"] =
get_predictions_of_model_on_data(classifier,
stemmed_bigram_tfidf_train,

```

```

stemmed_bigram_tfidf_test,
newsgroups_train.target)

    results_dict[name] = temp

```

## 6. Evaluate using multiple metrics

```

for method in results_dict:
    print(method)
    print("\naccuracy:")
    for preprocessing in results_dict[method]:
        print(preprocessing, accuracy_score(newsgroups_test.target,
            results_dict[method][preprocessing]))

    print("\nf1_score:")
    for preprocessing in results_dict[method]:
        print(preprocessing, f1_score(newsgroups_test.target,
            results_dict[method][preprocessing], average="weighted"))

    print("\nprecision:")
    for preprocessing in results_dict[method]:
        print(preprocessing, precision_score(newsgroups_test.target,
            results_dict[method][preprocessing], average="weighted"))

    print("\nrecall:")
    for preprocessing in results_dict[method]:
        print(preprocessing, recall_score(newsgroups_test.target,
            results_dict[method][preprocessing], average="weighted"))

```

### Naive Bayes

```

accuracy:
raw texts count 0.7728359001593202
raw texts tfidf 0.7738980350504514
preprocessed texts count 0.8147902283590016
preprocessed texts tfidf 0.8090812533191716
stemmed texts count 0.8035050451407328
stemmed texts tfidf 0.8036378120021243
stemmed bigram texts count 0.8098778544875199
stemmed bigram texts tfidf 0.8169144981412639

```

```

f1_score:
raw texts count 0.7511127577441177
raw texts tfidf 0.7684457156894653
preprocessed texts count 0.8091838550449055
preprocessed texts tfidf 0.8019006806034267
stemmed texts count 0.7952600019609147
stemmed texts tfidf 0.7956027121157477

```

stemmed bigram texts count 0.8047717132078135  
stemmed bigram texts tfidf 0.8104140121558414

precision:

raw texts count 0.7616683207318354  
raw texts tfidf 0.8218781741893993  
preprocessed texts count 0.8267221353663584  
preprocessed texts tfidf 0.8371801820381521  
stemmed texts count 0.8180376442726979  
stemmed texts tfidf 0.8342875141222853  
stemmed bigram texts count 0.8259120635209778  
stemmed bigram texts tfidf 0.8381756914860746

recall:

raw texts count 0.7728359001593202  
raw texts tfidf 0.7738980350504514  
preprocessed texts count 0.8147902283590016  
preprocessed texts tfidf 0.8090812533191716  
stemmed texts count 0.8035050451407328  
stemmed texts tfidf 0.8036378120021243  
stemmed bigram texts count 0.8098778544875199  
stemmed bigram texts tfidf 0.8169144981412639  
SVC

accuracy:

raw texts count 0.15108868826340946  
raw texts tfidf 0.8186404673393521  
preprocessed texts count 0.6812267657992565  
preprocessed texts tfidf 0.8203664365374402  
stemmed texts count 0.6949017525225704  
stemmed texts tfidf 0.8199681359532661  
stemmed bigram texts count 0.6865374402549124  
stemmed bigram texts tfidf 0.8230217737652682

f1\_score:

raw texts count 0.137243819157177  
raw texts tfidf 0.8190868537587032  
preprocessed texts count 0.6946980171540693  
preprocessed texts tfidf 0.821007569856839  
stemmed texts count 0.7004357724784672  
stemmed texts tfidf 0.8203890091464641  
stemmed bigram texts count 0.6951136956374109  
stemmed bigram texts tfidf 0.8232377081717388

precision:

raw texts count 0.42191640757596965  
raw texts tfidf 0.8293601659114187  
preprocessed texts count 0.7486584066498635  
preprocessed texts tfidf 0.8335437244436114  
stemmed texts count 0.7387123521930778



```
stemmed texts tfidf 0.8304145735728876
stemmed bigram texts count 0.7475758108149568
stemmed bigram texts tfidf 0.8368385566383851
```

recall:

```
raw texts count 0.15108868826340946
raw texts tfidf 0.8186404673393521
preprocessed texts count 0.6812267657992565
preprocessed texts tfidf 0.8203664365374402
stemmed texts count 0.6949017525225704
stemmed texts tfidf 0.8199681359532661
stemmed bigram texts count 0.6865374402549124
stemmed bigram texts tfidf 0.8230217737652682
Regression
```

accuracy:

```
raw texts count 0.7892989909718534
raw texts tfidf 0.8274030801911842
preprocessed texts count 0.7882368560807222
preprocessed texts tfidf 0.8266064790228359
stemmed texts count 0.7859798194370685
stemmed texts tfidf 0.8238183749336165
stemmed bigram texts count 0.8060276155071694
stemmed bigram texts tfidf 0.828996282527881
```

f1\_score:

```
raw texts count 0.7885050573077783
raw texts tfidf 0.8257218194818299
preprocessed texts count 0.7882743310555875
preprocessed texts tfidf 0.8251279298059621
stemmed texts count 0.785989255466463
stemmed texts tfidf 0.8224013186489373
stemmed bigram texts count 0.8056794561550035
stemmed bigram texts tfidf 0.8273396758486636
```

precision:

```
raw texts count 0.793171232182615
raw texts tfidf 0.8306575949769146
preprocessed texts count 0.7929613373013767
preprocessed texts tfidf 0.8307157256154403
stemmed texts count 0.7899530149476869
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/
_classification.py:1344: UndefinedMetricWarning: Precision is ill-
defined and being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

```
stemmed texts tfidf 0.8270476956033173
stemmed bigram texts count 0.8101888218083144
```

stemmed bigram texts tfidf 0.8352319135628984

recall:

raw texts count 0.7892989909718534

raw texts tfidf 0.8274030801911842

preprocessed texts count 0.7882368560807222

preprocessed texts tfidf 0.8266064790228359

stemmed texts count 0.7859798194370685

stemmed texts tfidf 0.8238183749336165

stemmed bigram texts count 0.8060276155071694

stemmed bigram texts tfidf 0.828996282527881

Random Forest

accuracy:

raw texts count 0.7634094530005311

raw texts tfidf 0.7632766861391397

preprocessed texts count 0.7646043547530537

preprocessed texts tfidf 0.7558417419012214

stemmed texts count 0.7579660116834838

stemmed texts tfidf 0.7590281465746149

stemmed bigram texts count 0.7738980350504514

stemmed bigram texts tfidf 0.7708443972384493

f1\_score:

raw texts count 0.7580805859262573

raw texts tfidf 0.7574100609173702

preprocessed texts count 0.7608021068324565

preprocessed texts tfidf 0.7508867634523942

stemmed texts count 0.7538081382423497

stemmed texts tfidf 0.754624747835952

stemmed bigram texts count 0.7712776722902094

stemmed bigram texts tfidf 0.7680483320123999

precision:

raw texts count 0.773647020129597

raw texts tfidf 0.7730665546564046

preprocessed texts count 0.774516481336767

preprocessed texts tfidf 0.7625820494144877

stemmed texts count 0.7683565620600574

stemmed texts tfidf 0.7695927148217405

stemmed bigram texts count 0.7908852618362442

stemmed bigram texts tfidf 0.7856574180761926

recall:

raw texts count 0.7634094530005311

raw texts tfidf 0.7632766861391397

preprocessed texts count 0.7646043547530537

preprocessed texts tfidf 0.7558417419012214

stemmed texts count 0.7579660116834838

stemmed texts tfidf 0.7590281465746149

```
stemmed bigram texts count 0.7738980350504514
stemmed bigram texts tfidf 0.7708443972384493
```

## 7. Questions

Which preprocessing approach and vector representation combination yields the best results in terms of classification accuracy?

The best result I got, after training all the selected models with both vector representations on different kinds of preprocessing is **0.828996282527881** in terms of accuracy.

I achieved it with **Logistic Regression trained on preprocessed texts including stemming, using tf-idf vectorizer that considers both bigrams and unigrams**

## 8. Calculate the F1-score using different values of the 'average' and 'labels'

```
best_model_predictions =
get_predictions_of_model_on_data(LogisticRegression(),
                                stemmed_bigram_tfidf_train,
                                stemmed_bigram_tfidf_test,
                                newsgroups_train.target)

print("accuracy of the model:", accuracy_score(newsgroups_test.target,
best_model_predictions))

print("average='micro', labels='None':",
f1_score(newsgroups_test.target, best_model_predictions,
average='micro', labels='None'))
print("average='micro':", f1_score(newsgroups_test.target,
best_model_predictions, average='micro'))

print("average='macro', labels='None':",
f1_score(newsgroups_test.target, best_model_predictions,
average='macro', labels='None'))
print("average='macro':", f1_score(newsgroups_test.target,
best_model_predictions, average='macro'))

print("average='weighted', labels='None':",
f1_score(newsgroups_test.target, best_model_predictions,
average='weighted', labels='None'))
print("average='weighted':", f1_score(newsgroups_test.target,
best_model_predictions, average='weighted'))

print("average='weighted', labels=[1,2,3,4,5]:",
f1_score(newsgroups_test.target, best_model_predictions,
average='weighted', labels=[1,2,3,4,5]))
```

accuracy of the model: 0.828996282527881

average='micro', labels='None': 0.7249999999999999

average='micro': 0.828996282527881

average='macro', labels='None': 0.5438401661406774

average='macro': 0.820492544745689

average='weighted', labels='None': 0.7253317849988162

average='weighted': 0.8273396758486636

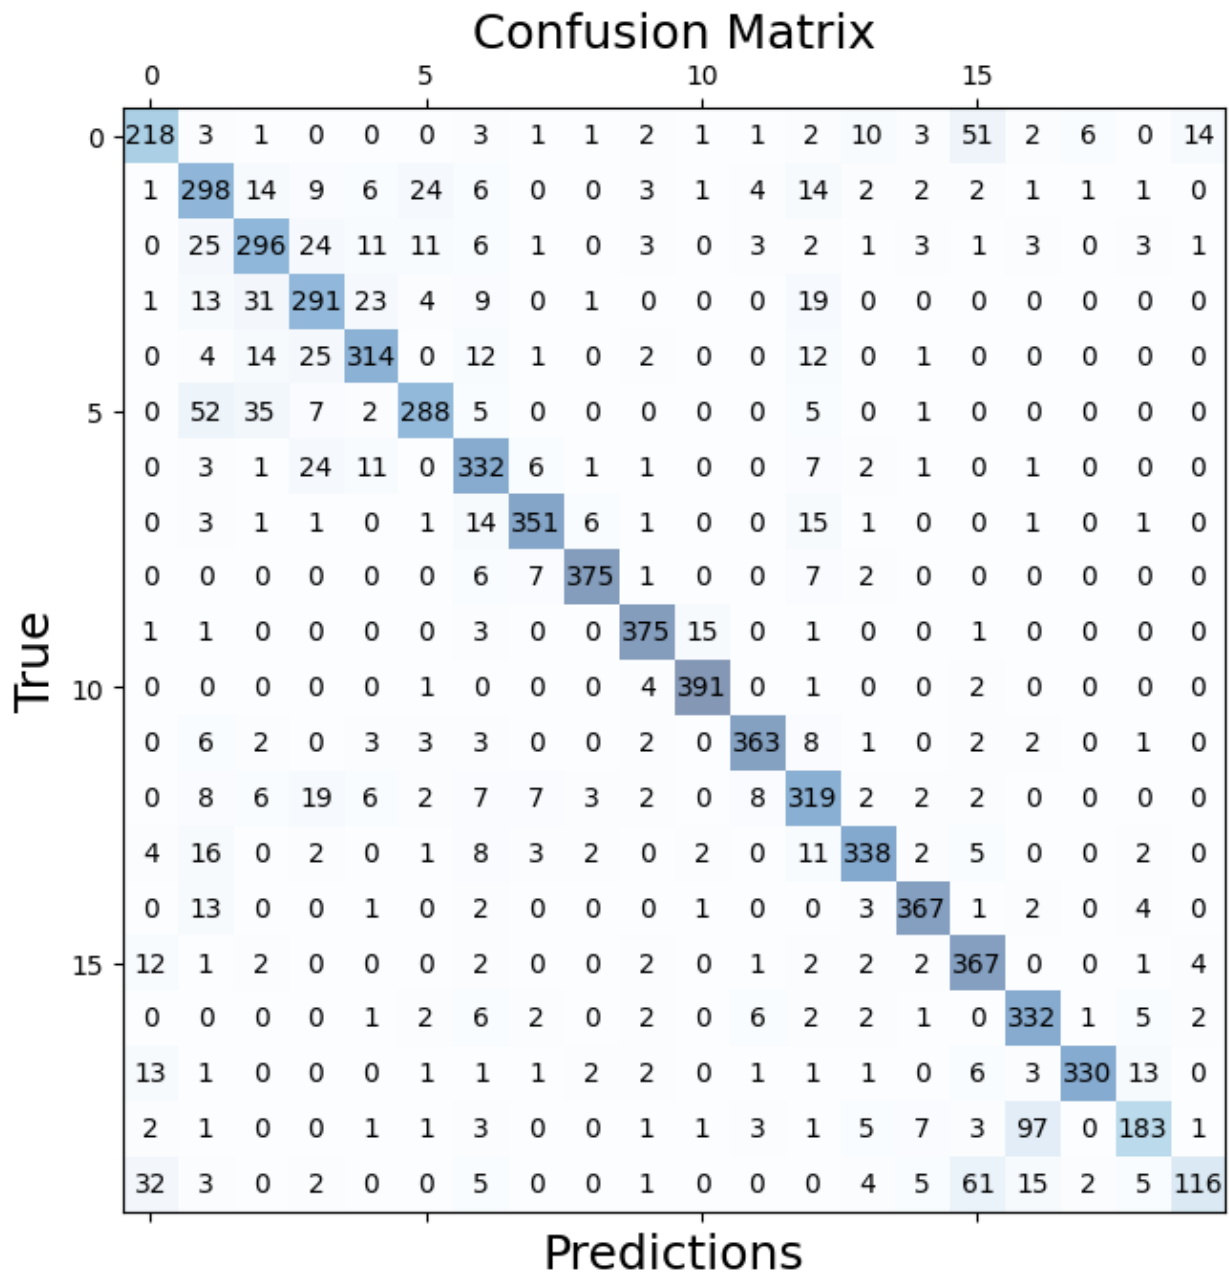
average='weighted', labels=[1,2,3,4,5]: 0.7579099812567712

## 9. Generate a confusion matrix

```
conf_matrix = confusion_matrix(newsgroups_test.target,
                                best_model_predictions)

fig, ax = plt.subplots(figsize=(7.5, 7.5))
ax.matshow(conf_matrix, cmap=plt.cm.Blues, alpha=0.5)
for i in range(conf_matrix.shape[0]):
    for j in range(conf_matrix.shape[1]):
        ax.text(x=j, y=i, s=conf_matrix[i, j], va='center',
                ha='center')

plt.xlabel('Predictions', fontsize=18)
plt.ylabel('True', fontsize=18)
plt.title('Confusion Matrix', fontsize=18)
plt.show()
```



## 10. Questions

What do micro, macro, or weighted values of the 'average' parameter of the F1-score allow you to set?

How the average is computed.

- Micro
  - Calculate metrics globally by counting the total true positives, false negatives and false positives.

- Macro
  - Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.
- weighted
  - Calculate metrics for each label, and find their average weighted by support (the number of true instances for each label). This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall.

from: [https://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1\\_score.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html)

What are the classes where the classification model tends to perform poorly or make the most mistakes?

As we can see from the confusion matrix, class 18 (`talk.politics.misc`) is being often classified as class 16 (`talk.politics.guns`), with 97 instances of class 18 (`talk.politics.misc`) ending up with the label 16 (`talk.politics.guns`).

Another problematic class is class 19 (`talk.religion.misc`), which is being classified as 15 (`soc.religion.christian`) with 61 mislabeled instances and as 0 (`alt.atheism`) with 32 mislabeled instances.

There is also quite a big confusion between classes 1 to 5, being the topics of `comp.graphics`, `comp.os.ms-windows.misc`, `comp.sys.ibm.pc.hardware`, `comp.sys.mac.hardware` and `comp.windows.x`.